

Running head: SURVIVAL AND LATENCIES

Are divergence point analyses suitable ~~to examine~~ latency data?

Pablo Gomez (1), Javier Breithaupt (2), Manuel Perea(2, 3), Jeff (4)

(1) DePaul University, Chicago, USA

(2) Universitat de València, València, Spain

(3) BCBL, Basque Center on Cognition, Brain, and Language, San Sebastián, Spain

(4) University of Missouri

Pablo Gomez

Department of Psychology

DePaul University

Chicago, Illinois

60614

e-mail: pgomez1@depaul.edu

Abstract

The uncovering of the time course of the influence of different factors in human performance is one of the principal topics of research in cognitive psychology/neuroscience. Over the past decades, researchers have proposed several methods to tackle this question using latency data. Here we examined a recently proposed procedure that employs survival analyses on latency data to provide “precise estimates” of the timing of the first discernible influence of a given factor on performance (e.g., word frequency on lexical access). A number of articles have used this method in recent years, and hence an exploration of its strengths and its potential weaknesses is in order. Unfortunately our analysis revealed that the method has fundamental conceptual flaws, and it might lead researchers into believing that they are obtaining a measurement of a processing component when in fact they are obtaining a non-sensical measurement.

Are divergence point analyses suitable to examine latency data?

Perhaps the most common cognitive psychology experiment is one in which participants are presented with ~~a stimulus that has some feature~~ of theoretical interest (e.g., length, word frequency, etc.). The stimulus elicits a response, and researchers often measure latencies to make inferences about hypothesized underlying cognitive processes. This form of mental chronometry is widely used in the analyses of data from a wide range of experimental paradigms such as choice tasks, naming, eye-tracking, and many others.

~~Although most often the obtained effects are examined through statistical inference methods on the mean latencies, the shortcomings of focusing only on the mean latencies~~ are well known (Balota & Yap, 2011; Heathcote, Popiel, & Mewhort, 1991; Ratcliff, 1979); hence, researchers benefit from exploring distributional properties of latency measurements. **Some are analyses** based on fitting functional forms like the ex-Gaussian or the Weibull distributions (see Heathcote et al., 1991), and other methods are based on process models like the diffusion model for choice response times (Ratcliff, 1978) or the EZ-reader model for eye fixation durations during reading (Reichle, Pollatsek, Fisher, & Rayner, 1998).

Is there a method that goes beyond mean latency analyses without making strong theoretical commitments or specific assumptions about distributional properties of the latency measurements and does not require a specific experimental paradigm? In this note, we ~~explore a method that might~~ fulfill that void in the literature, namely, the *divergence point analysis* (Reingold, Reichle, Glaholt, & Sheridan, 2012; Sheridan, 2013). **Its goal** is to provide an estimate of the onset of the influence of a given variable (the divergence point) on the basis of latency data (e.g., response times or eye fixation times).



the context of latency measurements, if T is a random variable with **CDF** $F(t)$;

its survival function on time (t) is defined as $S(t) = P(T > t) = 1 - F(t)$; where the survival at time t is equal to the proportion of latencies that have not occurred by time t ; hence at $t = 0$, $S(0) = 1$ and at $t = \infty$, $S(\infty) = 0$. There have been attempts to use survival functions in latency analyses. Notably, Van Zandt (2002) examined several of these procedures and concluded that serious analyses of this type “would use samples of at least a few hundred observations” (p. 482), along similar lines Houpt and Townsend (2010, 2011) discussed a rather sophisticated method of survival analysis that involves experimental methods and a non-trivial algorithm termed survivor interaction contrasts (SIC).

Recently, Sheridan, Reingold, and colleagues (Sheridan, 2013; see also Ando, Matsuki, Sheridan, & Jared, 2015; Reingold & Sheridan, 2014; Reingold, et al., 2012; Sheridan & Reingold, 2013; Sheridan, Rayner & Reingold, 2013) proposed a method that promises to overcome the limitations of traditional survival function analyses by using a computationally intensive bootstrapping procedure. This procedure compares the latency distributions of two conditions by estimating the divergence point. The divergence point corresponds to the shortest latency value at which a manipulation has a significant impact. Thus, the divergence point offers an estimate of the timing of the first discernible influence of a given variable (e.g., word frequency in a word recognition task). Furthermore, this method can also inform us whether the effect of a given factor has an earlier onset than the effect of another factor.

The basic setup in the *divergence point* method under consideration can be described as follows: In each iteration the latencies for each participant and condition are randomly re-sampled with replacement. These sampled (bootstrapped) latencies are used to compute each participant’s survival curves, which in turn are averaged across subjects *à la* Vincentile (Vincent, 1912). Next, for each time bin t the difference between conditions: $\Delta_{t,i}$ are computed (i^{th} iteration from 1 to 10,000), and then sorted. The range

between the 5th and the 9,995th value becomes the confidence interval $CI(\Delta t)$ and the divergence point is defined as the shortest t at which the $CI(\Delta t)$ does not include 0. Aiming for a high temporal resolution, Sheridan and Reingold use 1ms bins (see Figure 1 for an example).

Although this method might seem promising and useful for researchers interested in exploring the time course of a given empirical effect, there is a fundamental conceptual flaw in the foundation of the method.



Estimating divergence points in latency distributions is conceptually flawed

The divergence point method is a way to estimate the point at which the latency distributions corresponding to two different experimental conditions separate. The underlying assumption behind the method is that there might be experimental situations in which two conditions share a process, and this is followed by a subsequent processing component in which the two conditions differ (diverge) on the timing and/or the processing cost. Although one could think of somewhat contrived situations in which this type of serial, staged process might occur (see Carreiras, Armstrong, Perea, & Frost, 2014, for a recent review against this type of formulation), the latencies, even at their shortest duration, would be a reflection of the **ending times** of the second process. In other words, the underlying divergence is not reflected by the latency. Latency distributions reflect aggregation across processes, across trials, and even across participants.

The interpretation of the output provided by this method can be misleading, as it provides nonsensical answers when applied to real data as will be clear with a few examples: Figures 2 and 3 display the cumulative density functions generated with an ex-Gaussian distribution in which there are effects on μ (Figure 2) and on τ (Figure 3). The interpretation of both figures is quite straightforward: the distributions are separated

from the very onset of the CDFs. The divergence point is at the starting point of the latency distribution any time there is *stochastic dominance*, as is the case for the overwhelming majority of experimental manipulations in psychology. Stochastic dominance refers to the probability of observations smaller than x being greater for one variable than the other for all values of x (see Heathcote, Brown, Wagenmakers, & Eidels, 2010).

The method provides non-sensical results

What are the consequences of applying this method to data in which there is stochastic dominance? To answer this question, we carried out a series of Monte Carlo simulations. In particular, we generated data from an ex-Gaussian distribution assuming that the experimental effect was either in the μ or τ parameters. We manipulated the number of hypothetical trials per condition and then applied the bootstrapping method to estimate the divergence point.

For the first simulation, we explored if the method yields false positives when samples are generated from an ex-Gaussian with parameters $\mu = 541$, $\sigma = 68$, and $\tau = 115$ (i.e., a null effect in which identical parameters generate the simulated data for the two conditions; note that these parameters were the average parameters in Heathcote et al, 1991). The results from this simulation were satisfactory: The method generates false divergence points less than .01% regardless of the number of items in the simulations.

In the second simulation, we generated data in which the difference between the two conditions was an effect on μ . The data for the baseline condition was generated from an ex-Gaussian distribution with $\mu = 541$, $\sigma = 68$, and $\tau = 115$. We generated data for three simulated experimental conditions by changing μ to 561, 581 and 621 ($\Delta\mu = 20, 40, 80$). There is, therefore, stochastic dominance of the baseline condition relatively to all of these other conditions, and the true divergence point is at the starting point of the

distributions. The results from this simulation are not encouraging for the method, as the estimation of the divergence point is highly biased by the number of trials per condition ($n = 20, 30, 50, 100, 250, 500, 1000$). Figure 5 shows the average divergence point for each of the parameter combinations (μ and n) across 1000 simulations. As a consequence of increased statistical power due to larger sample size at the trial level, the larger the number of trials per condition, the shorter the estimated divergence point (i.e., there is a statistical bias dependent on sample size). For example, with $\Delta\mu = 80$, and an n of 50, the divergence point is about 100ms higher than for $n = 1000$. In fact, when the number of trials is below 100, the different conditions are undistinguishable from each other. While the use of confidence intervals proposed by Reingold and Sheridan (2014) might partially alleviate this problem, the existence of a bias dependent on sample size is inherent to the procedure.

In the third simulation, we generated latency data in which the difference between the two conditions was an effect on τ . The data for the baseline condition was the same as in the previous simulations: it was generated from an ex-Gaussian distribution with $\mu = 541$, $\sigma = 68$, and $\tau = 115$. We generated data for three simulated experimental conditions by changing τ to 135, 155 and 195 ($\Delta\tau = 20, 40, 80$). As shown in Figure 2, changes in τ produce stochastic dominance of the baseline condition and the true divergence point is at the starting point of the distributions. The results from this simulation are very similar to those from Simulation 2: The estimation of the divergence point is severely biased by the number of trials per condition ($n = 20, 30, 50, 100, 250$ or 500). Figure 6 shows the average divergence point for each of the parameter combinations (τ and n) across 1000 simulations.

In sum, the method of estimating divergence points from two latency distributions does not generate false positives beyond the α level stated by the method; but more importantly, when two latency distributions differ, the method provides an answer that is

heavily determined by the number of observations per condition.

Conclusion

Although the goals of the divergence point method are worth pursuing, our analysis revealed serious shortcomings on the conceptual foundation of the procedure: Latency measurements tend to exhibit stochastic dominance between experimental conditions, and hence the divergence point would be at the leading edge of the latency distribution regardless of other distributional differences. Furthermore, if the method is applied to data, an estimate of the divergence point will be provided by the method. This estimate will be affected mostly by the number of observations. In short, our exploration of the method forces us to conclude that it is not advisable to utilize it when analyzing latency data.

References

- Ando, E., Matsuki, K., Sheridan, H., & Jared, D. (2015). The locus of Katakana-English masked phonological priming effects. *Bilingualism: Language and Cognition*, 18, 101-117. DOI: [dx.doi.org/10.1017/S1366728914000121](https://doi.org/10.1017/S1366728914000121)
- Balota, D. A., & Yap, M. J. (2011). Moving beyond the mean in studies of mental chronometry the power of response time distributional analyses. *Current Directions in Psychological Science*, 20, 160-166. doi: [10.1177/0963721411408885](https://doi.org/10.1177/0963721411408885)
- Carreiras, M., Armstrong, B.C., Perea, M. & Frost, R. , (2014) The What, When, Where, and How of Visual Word Recognition. *Trends in Cognitive Sciences*, 18, 90-98.
- Gomez, P., Perea, M., & Ratcliff, R. (2013). A diffusion model account of masked vs. unmasked priming: Are they qualitatively different? *Journal of Experimental Psychology: Human Perception and Performance*, 39, 1731-1740. DOI: [10.1037/a0032333](https://doi.org/10.1037/a0032333)
- Heathcote, A., Popiel, S. J., & Mewhort, D. J. (1991). Analysis of response time

distributions: An example using the Stroop task. *Psychological Bulletin*, 109, 340. DOI: 10.1037/0033-2909.109.2.340

Heathcote, A., Brown, S., Wagenmakers, E. J., & Eidels, A. (2010).

Distribution-free tests of stochastic dominance for small samples. *Journal of Mathematical Psychology*, 54, 454-463. doi:10.1016/j.jmp.2010.06.005

Ratcliff, R. (1978). A theory of memory retrieval. *Psychological Review*, 85, 59-108. doi:10.1037/0033-295X.85.2.59

Ratcliff, R. (1979). Group reaction time distributions and an analysis of distribution statistics. *Psychological Bulletin*, 86, 446-461. doi: dx.doi.org/10.1037/0033-2909.86.3.446

Reichle, E. D., Pollatsek, A., Fisher, D. L., & Rayner, K. (1998). Toward a model of eye movement control in reading. *Psychological Review*, 105, 125 -157.

Reingold, E. M., Reichle, E. D., Glaholt, M. G., & Sheridan, H. (2012). Direct lexical control of eye movements in reading: Evidence from a survival analysis of fixation durations. *Cognitive Psychology*, 65, 177-206.

Reingold, E. M., & Sheridan, H. (2014). Estimating the divergence point: a novel distributional analysis procedure for determining the onset of the influence of experimental variables. *Frontiers in Psychology*, 5, 1432. doi: 10.3389/fpsyg.2014.01432

Sheridan, H. (2013). The Time-course of Lexical Influences on Fixation Durations during Reading: Evidence from Distributional Analyses. Doctoral Dissertation, University of Toronto. Retrieved from https://tspace.library.utoronto.ca/bitstream/1807/35996/6/Sheridan_Heathe_201306_PhD_thesis.pdf

Sheridan, H., & Reingold, E. M. (2013). Distinct stages of word identification during reading: Evidence from eye movements. *Journal of Vision*, 13, 511-511.

Sheridan, H., Rayner, K., & Reingold, E. M. (2013). Unsegmented text delays word identification: Evidence from a survival analysis of fixation durations. *Visual Cognition*, 21, 38-60.

Houpt, J. W., & Townsend J. T. (2010). The statistical properties of the survivor interaction contrast. *Journal of Mathematical Psychology*, *54*, 446-453.

doi:10.1016/j.jmp.2010.06.006

Houpt, J. W., & Townsend, J. T. (2011). An extension of SIC predictions to the Wiener coactive model, *Journal of Mathematical Psychology*, *55*, 267-270. doi:

dx.doi.org/10.1016/j.jmp.2011.02.002.

Van Zandt, T. (2002). Analysis of response time distributions. In J. T. Wixted (Ed.), *Stevens' Handbook of Experimental Psychology (3rd ed., pp. 461-516)*. San Diego, CA: Academic Press.

Vincent, S. B. (1912). The function of vibrissae in the behavior of the white rat. *Behavioral Monographs*, *1* (Whole No. 5).

Author Note

The research reported in this article has been partially supported by Grant PSI2014-53444-P from the Spanish Ministry of Economy and Competitiveness. Pablo Gomez was the recipient of the “Convitat” grant at the University of Valencia (VLC-Campus).

Figure Captions

Figure 1. The figure (modified from Sheridan, 2013, page 27) shows the distributions of first-fixation duration on target words in the low and high predictability conditions in the top panel, and the survival curves in the bottom panel. The row of points at the top of the survival curves indicates the time bins with a significant difference between the low and high predictability curves using the method being examined in the present note

Figure 2. The figure shows the cumulative density functions generated with an ex-Gaussian distribution in which there are effects on μ

Figure 3. The figure shows the cumulative density functions generated with an ex-Gaussian distribution in which there are effects on τ

Figure 4. The figure shows average diverge point for simulation 2.

Figure 5. The figure shows average diverge point for simulation 3.









